

Target Papers Shortlist

20 target papers ranked by composite relevance, each with the validated follow-up studies it has prompted.

Probing (7)Steering (13)

STUDY		RELEVANCE		LINKS	
1	AxBench: Steering LLMs? Even Simple Baselines Outperform Sp... Steering · Gemma · 2–9B · 3 extensions · ICML 2025	63.5		ICML	
	SAEs Are Good for Steering -- If You Select the Right Features · EMNLP 2025			arXiv	
	Does higher interpretability imply better utility? A Pairwise Analysis on Sparse Autoencoders			arXiv	
	Enhancing LLM Steering through Sparse Autoencoder-Based Vector Refinement			arXiv	
2	DISCO: Disentangled Communication Steering for Large Langua... Steering · Gemma, Llama · 8–9B · NeurIPS 2025	60.2		GitHub OpenReview	
3	Detecting Strategic Deception with Linear Probes Probing · Llama · 70B · 1 extension · ICML 2025	58.2		GitHub ICML	
	The Secret Agenda: LLMs Strategically Lie and Our Current Safety Tools Are Blind · AAAI 2026 Workshop			arXiv	
4	Detecting High-Stakes Interactions with Activation Probes Probing · NeurIPS 2025	57.5		GitHub OpenReview	
5	Improving Instruction-Following in Language Models through... Steering · 4 extensions · ICLR 2025	55.0		GitHub ICLR	
	Enhancing Instruction Following of LLMs via Activation Steering with Dynamic Rejection			arXiv	
	Prompt-Activation Duality: Improving Activation Steering via Attention-Level Interventions			arXiv	
	Steering Code LLMs with Activation Directions for Language and Library Control			arXiv	
	Denoising Concept Vectors with Sparse Autoencoders for Improved Language Model Steering			arXiv	
6	Angular Steering: Behavior Control via Rotation in Activati... Steering · NeurIPS 2025	54.2		GitHub OpenReview	
7	Programming Refusal with Conditional Activation Steering Steering · 1 extension · ICLR 2025	54.0		ICLR	
	C-ΔO: Circuit-Restricted Weight Arithmetic for Selective Refusal			arXiv	
8	SCANS: Mitigating the Exaggerated Safety for LLMs via Safet... Steering · AAAI 2025	53.3		GitHub AAAI	
9	LayerNavigator: Finding Promising Intervention Layers for E... Steering · NeurIPS 2025	53.0		GitHub OpenReview	
10	Towards Interpretability Without Sacrifice: Faithful Dense... Probing · 3B · NeurIPS 2025	52.7		GitHub OpenReview	
11	REAL: Reading Out Transformer Activations for Precise Local... Steering · Llama, Qwen · ICLR 2026	51.0		OpenReview	
12	The Linear Representation Hypothesis and the Geometry of La... Probing · Llama · 3 extensions · ICML 2024	50.0		ICML	
	The Geometry of Categorical and Hierarchical Concepts in Large Language Models · ICLR 2025 (oral)			arXiv	
	Toward a Flexible Framework for Linear Representation Hypothesis Using Maximum Likelihood Estimation			arXiv	
	Mapping Semantic & Syntactic Relationships with Geometric Rotation			arXiv	
13	InferAligner: Inference-Time Alignment for Harmlessness thr... Steering · 1 extension · EMNLP 2024	49.7		GitHub ACL	
	Inference-time Alignment in Continuous Space			arXiv	
14	Embedded Named Entity Recognition using Probing Classifiers Probing · EMNLP 2024	49.0		GitHub ACL	
15	Steering Knowledge Selection Behaviours in LLMs via SAE-Bas... Steering · NAACL 2025	48.4		GitHub ACL	
16	Activation Steering for LLM Alignment via a Unified ODE-Bas... Steering · ICLR 2026	48.3		OpenReview	
17	Differentially Private Steering for Large Language Model Al... Steering · Llama, Qwen · 0.5–7B · ICLR 2025	46.7		ICLR	
18	LEGEND: Leveraging Representation Engineering to Annotate S... Steering · AAAI 2025	46.0		GitHub AAAI	
19	Beyond Linear Probes: Dynamic Safety Monitoring for Languag... Probing · 30B · ICLR 2026	46.0		OpenReview	
20	Are Sparse Autoencoders Useful? A Case Study in Sparse Prob... Probing · 2 extensions · ICML 2025	44.3		ICML	
	Rethinking Sparse Autoencoders: Select-and-Project for Fairness and Control from Encoder Features Alone			arXiv	
	Inference-Time Decomposition of Activations (ITDA): A Scalable Approach to Interpreting Large Language Models			arXiv	

- Relevance
- Composite of focus, reproducibility, cross-model readiness, impact and recency.
- Extension
- A validated follow-up study; hover any row for the linking evidence.
- Links
- Buttons open the GitHub repo and the source page; each extension opens its own link.