Six of eight original-model baselines reproduce

Deviation of our rerun from the published value, on the paper's own model.

Every percentage-point metric shares one scale.

● Baseline recovered ● Baseline not recovered ▭ Within 1 pp of published

Percentage-point metrics, shared scale

	Published	Our rerun	Deviation
Models Under Pressure — Mean AUROC, six datasets	91.70	**91.84**	+0.14
Instruction Following — Mean AUROC, held-out cells			
Unseen tasks	68.44	**68.46**	+0.02
Unseen types	52.17	**52.01**	−0.16
Lookback Lens — Mean cross-task transfer AUROC	74.85	**74.80**	−0.05
Beyond Linear Probes — WildGuardMix test F1, selected layer			
Linear	83.24	**83.26**	+0.02
Bilinear	84.78	**84.35**	−0.43
MLP	83.77	**83.66**	−0.11
EE-MLP	83.84	**83.92**	+0.08
TPC	84.48	**84.81**	+0.33
PRISM — Prompted minus raw accuracy			
MM	18.91	**18.91**	+0.00
SAPLMA	13.42	**14.13**	+0.71
Among Us — Mean AUROC, 20 cells	84.68	**73.39**	−11.29
Sycophancy — Peak validation accuracy			
MLP	98.09	**79.13**	−18.96
Residual stream	98.47	**79.26**	−19.21

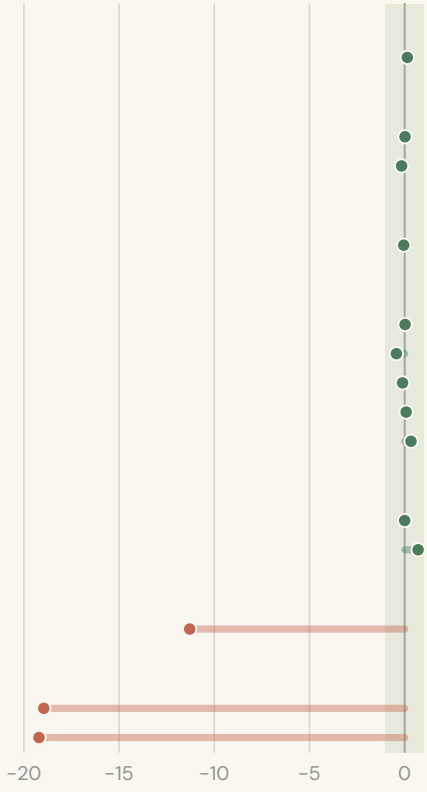

Correlation metric, its own scale

	Published	Our rerun	Deviation
Political Perspective — Peak single-head Spearman correlation			
Llama-2	0.8536	**0.8536**	+0.00001
Mistral	0.8463	**0.8462**	−0.00013
Vicuna	0.8612	**0.8611**	−0.00011

Positive deviation means our rerun scored above the published value; the shaded band marks the ±1 pp recovery tolerance. Sycophancy is a five-seed mean; every other row is a single estimate.